

Searching: Deterministic single-agent

Andrew W. Moore
Professor
School of Computer Science
Carnegie Mellon University

www.cs.cmu.edu/~awm

awm@cs.cmu.edu


412-268-7599

Note to other teachers and users of these slides. Andrew would be delighted if you found this source material useful in giving your own lectures. Feel free to use these slides verbatim, or to modify them to fit your own needs. PowerPoint originals are available. If you make use of a significant portion of these slides in your own lecture, please include this message, or the following link to the source repository of Andrew's tutorials: <http://www.cs.cmu.edu/~awm/tutorials> . Comments and corrections gratefully received.

Overview

- Deterministic, single-agent, search problems
- Breadth First Search
- Optimality, Completeness, Time and Space complexity
- Search Trees
- Depth First Search
- Iterative Deepening
- Best First “Greedy” Search

A search problem



How do we get from S to G? And what's the smallest possible number of transitions?


Formalizing a search problem

A search problem has five components:

Q , S , G , **succs** , **cost**

- Q is a finite set of states.
- $S \subseteq Q$ is a non-empty set of start states.
- $G \subseteq Q$ is a non-empty set of goal states.
- **succs** : $Q \rightarrow P(Q)$ is a function which takes a state as input and returns a set of states as output. **succs**(s) means “the set of states you can reach from s in one step”.
- **cost** : $Q \times Q \rightarrow \text{Positive Number}$ is a function which takes two states, s and s' , as input. It returns the one-step cost of traveling from s to s' . The cost function is only defined when s' is a successor state of s .

Our Search Problem



$Q = \{ \text{START}, a, b, c, d, e, f, h, p, q, r, \text{GOAL} \}$

$S = \{ \text{START} \}$

$G = \{ \text{GOAL} \}$


$\text{succs}(b) = \{ a \}$

$\text{succs}(e) = \{ h, r \}$

$\text{succs}(a) = \text{NULL} \dots \text{etc.}$

$\text{cost}(s, s') = 1$ for all transitions

Our Search Problem



$Q = \{ \text{START}, a, b, c, d, e, f, h, p, q, r, \text{GOAL} \}$

$S = \{ \text{START} \}$

$G = \{ \text{GOAL} \}$

$\text{succs}(b) = \{ a \}$

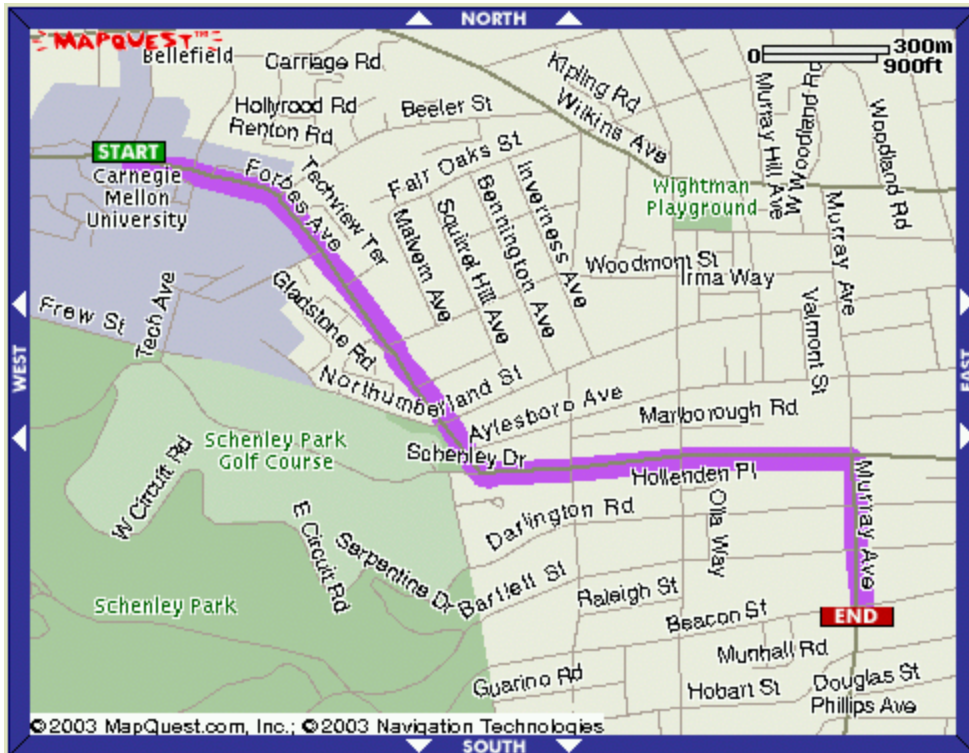

$\text{succs}(e) = \{ h, r \}$

$\text{succs}(a) = \text{NULL} \dots \text{etc.}$

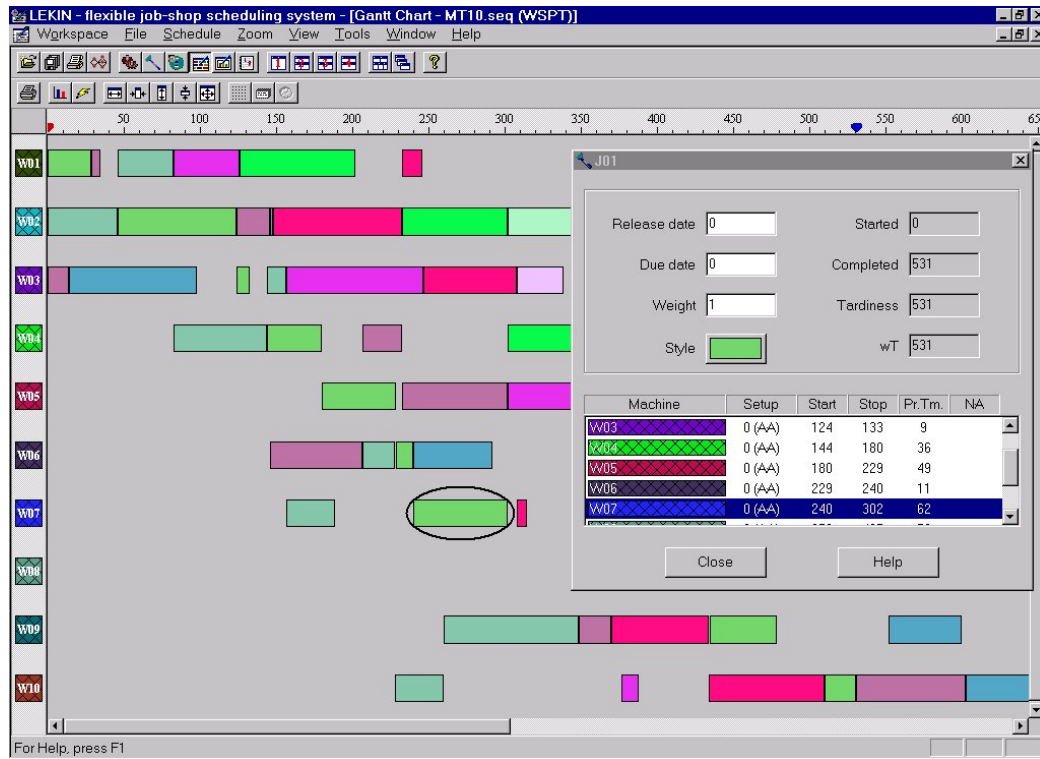
$\text{cost}(s, s') = 1$ for all transitions

Why do we care? What problems are like this?

Search Problems




More Search Problems




Scheduling

8-Queens



What

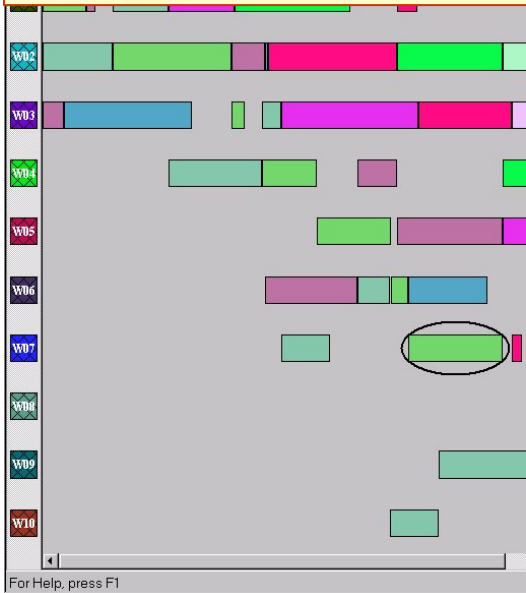


More Search Problems

But there are plenty of things which we'd normally call search problems that don't fit our rigid definition...

including


- A search problem has five components:
- Q , S , G , **succs**, **cost**
- Q is a finite set of states.
- $S \subseteq Q$ is a non-empty set of start states.
- $G \subseteq Q$ is a non-empty set of goal states.
- **succs** : $Q \rightarrow P(Q)$ is a function which takes a state as input and returns a set of states as output. **succs**(s) means "the set of states you can reach from s in one step".
- **cost** : $Q, Q \rightarrow \text{Positive Number}$ is a function which takes two states, s and s' , as input. It returns the one-step cost of traveling from s to s' . The cost function is only defined when s' is a successor state of s .



Can you think of examples?



Our definition excludes...



Our definition excludes

Game
against
adversary




Chance



Hidden State




Continuum (infinite
number) of states



All of the above, plus
distributed team control



Breadth First Search




Label all states that are reachable from S in 1 step but aren't reachable in less than 1 step.

Then label all states that are reachable from S in 2 steps but aren't reachable in less than 2 steps.

Then label all states that are reachable from S in 3 steps but aren't reachable in less than 3 steps.

Etc... until Goal state reached.


Breadth-first Search



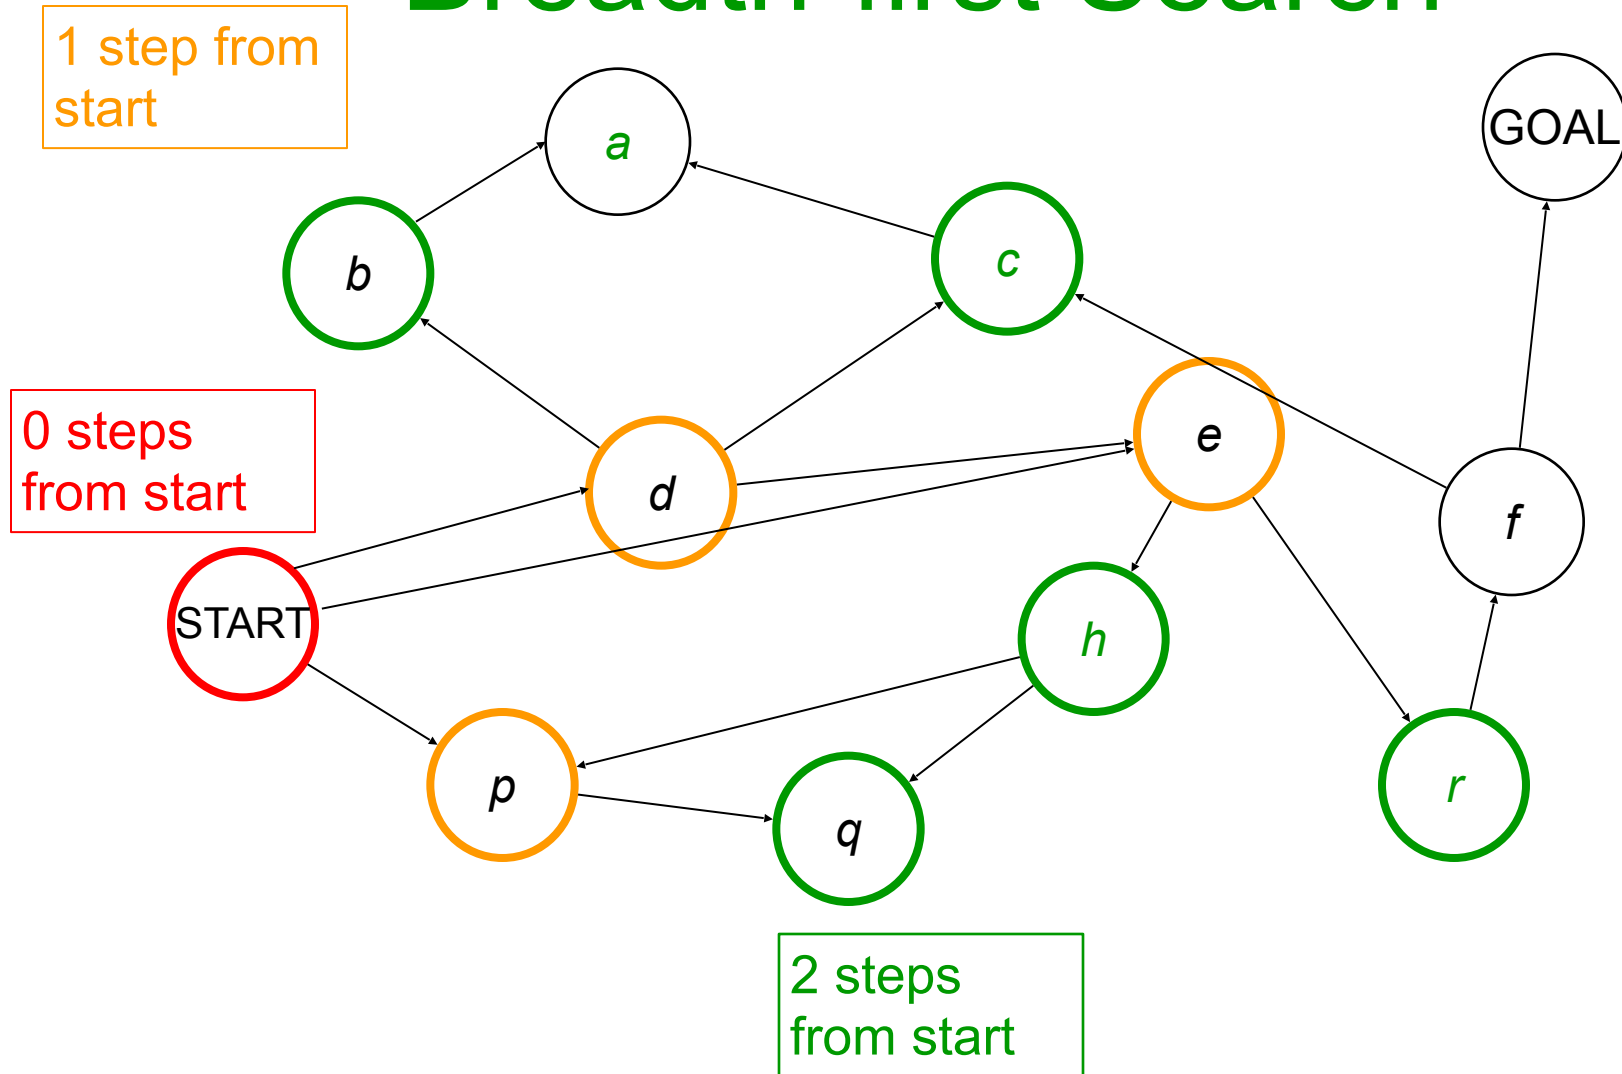
Breadth-first Search

1 step from start

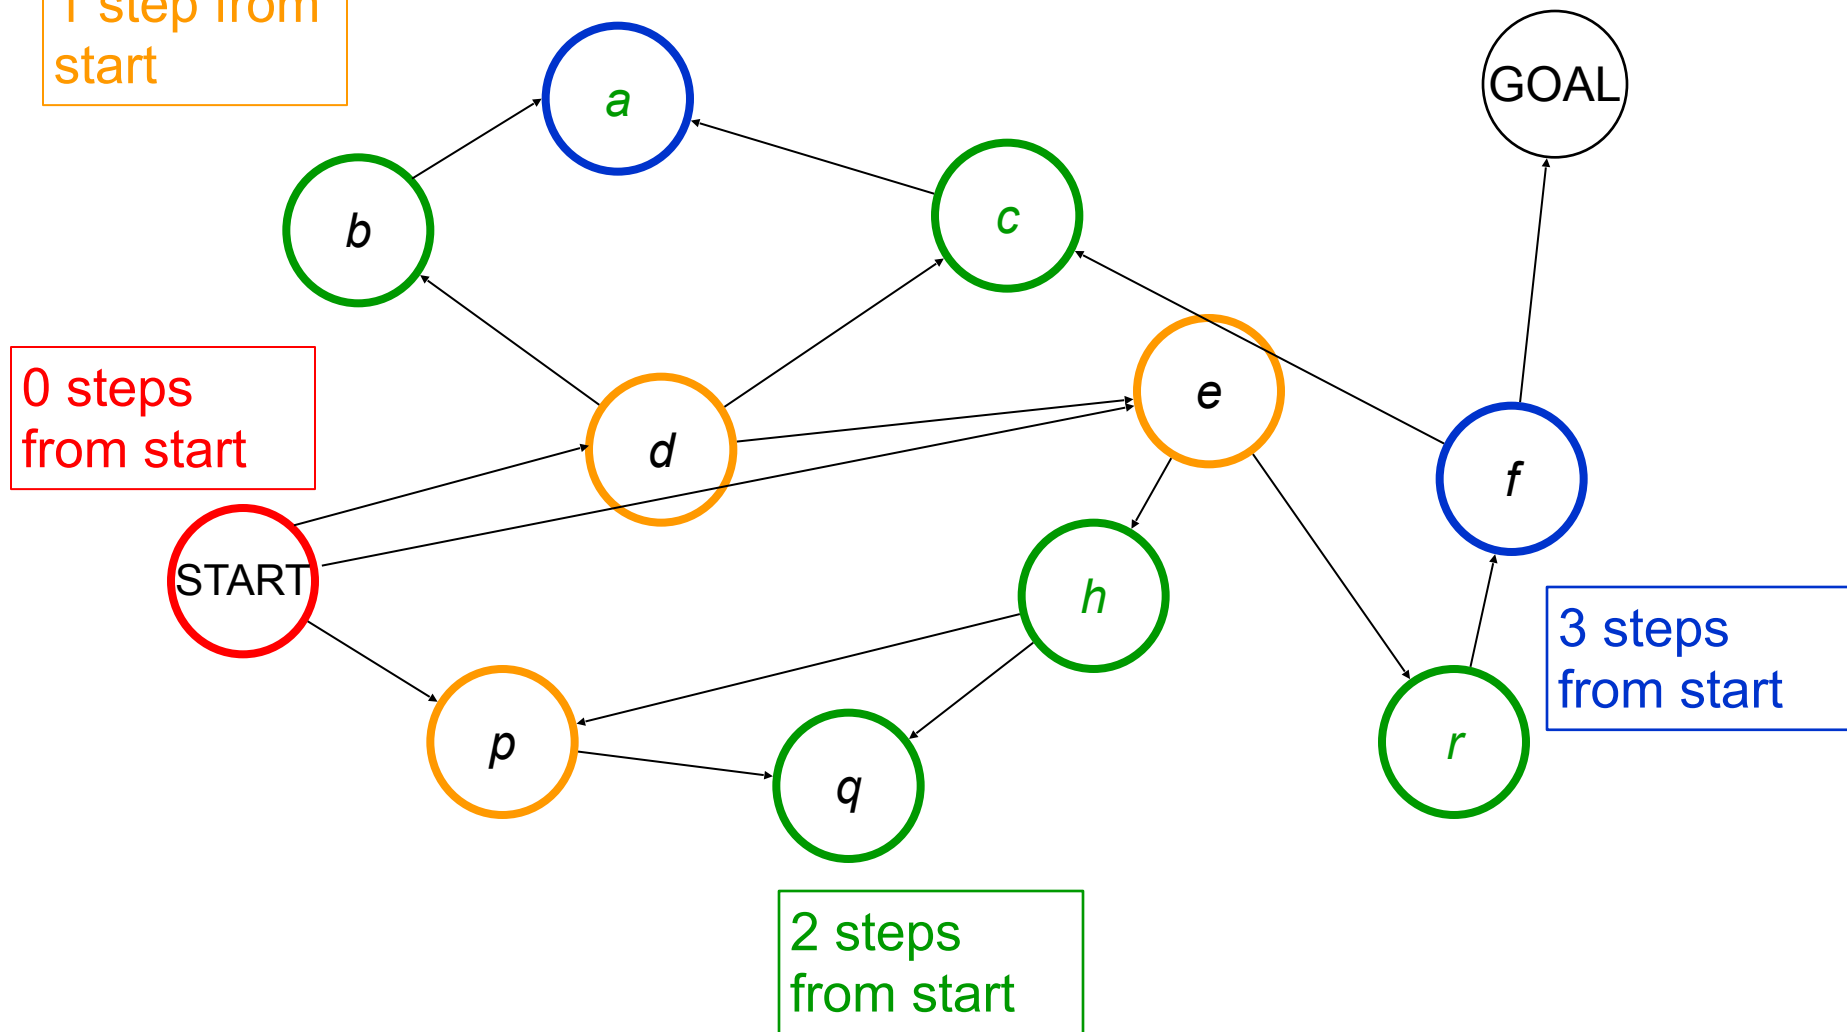
0 steps from start



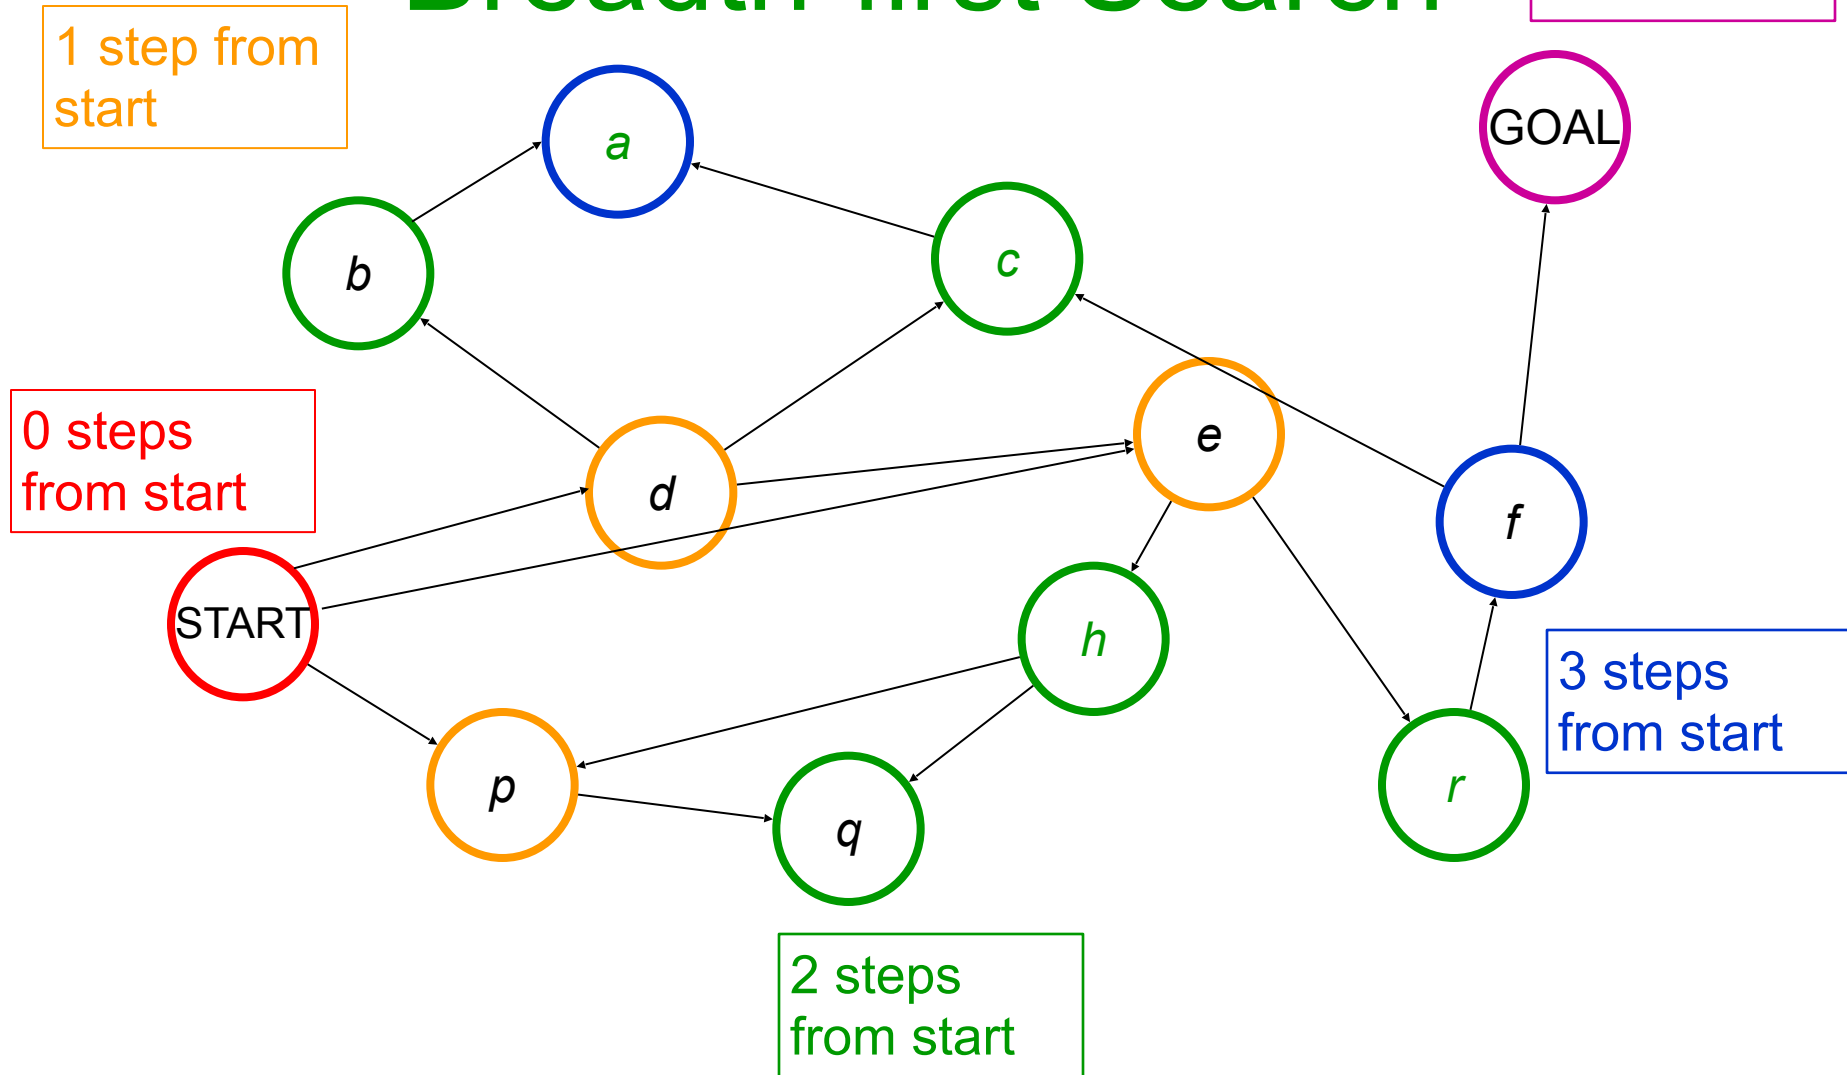
Breadth-first Search




Breadth-first Search



Breadth-first Search



Remember the path!

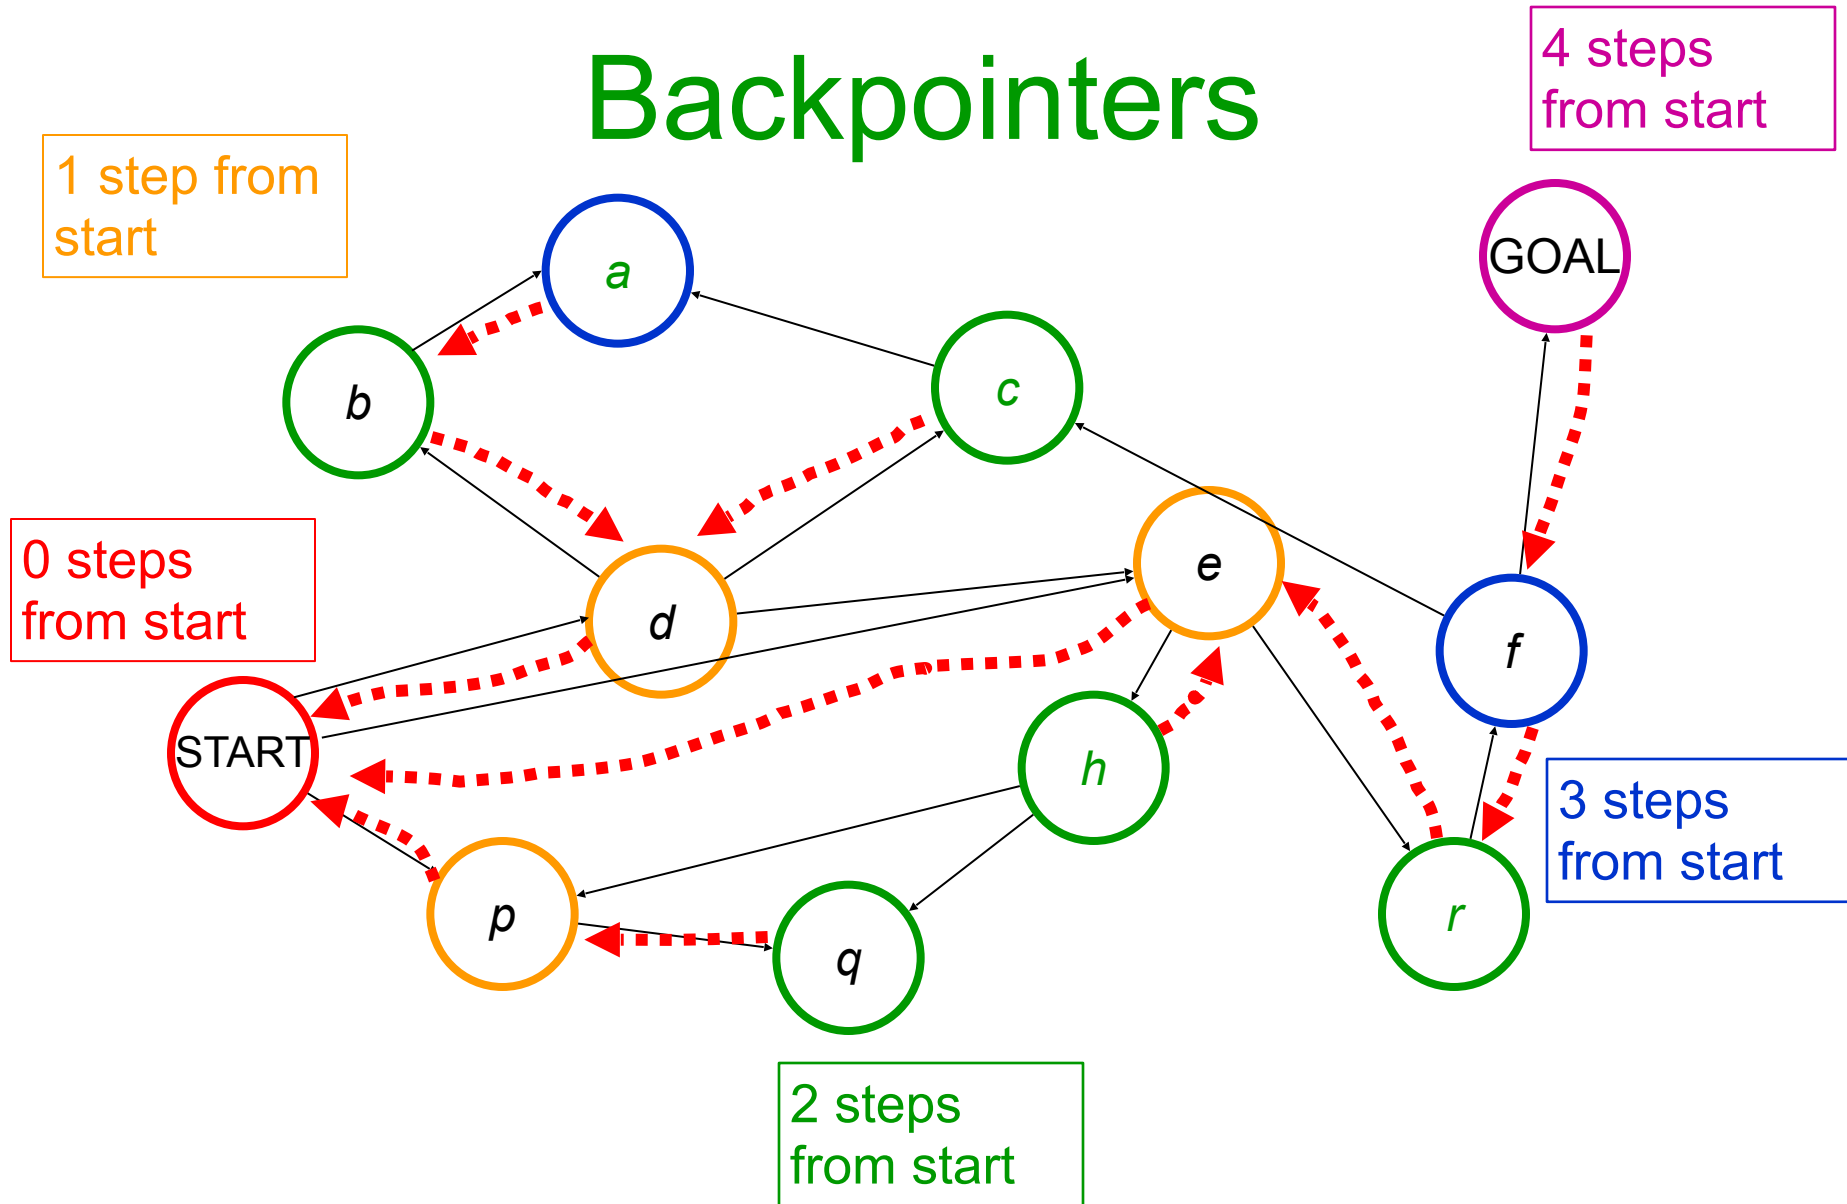


Also, when you label a state, record the predecessor state. This record is called a *backpointer*. The history of predecessors is used to generate the solution path, once you've found the goal:

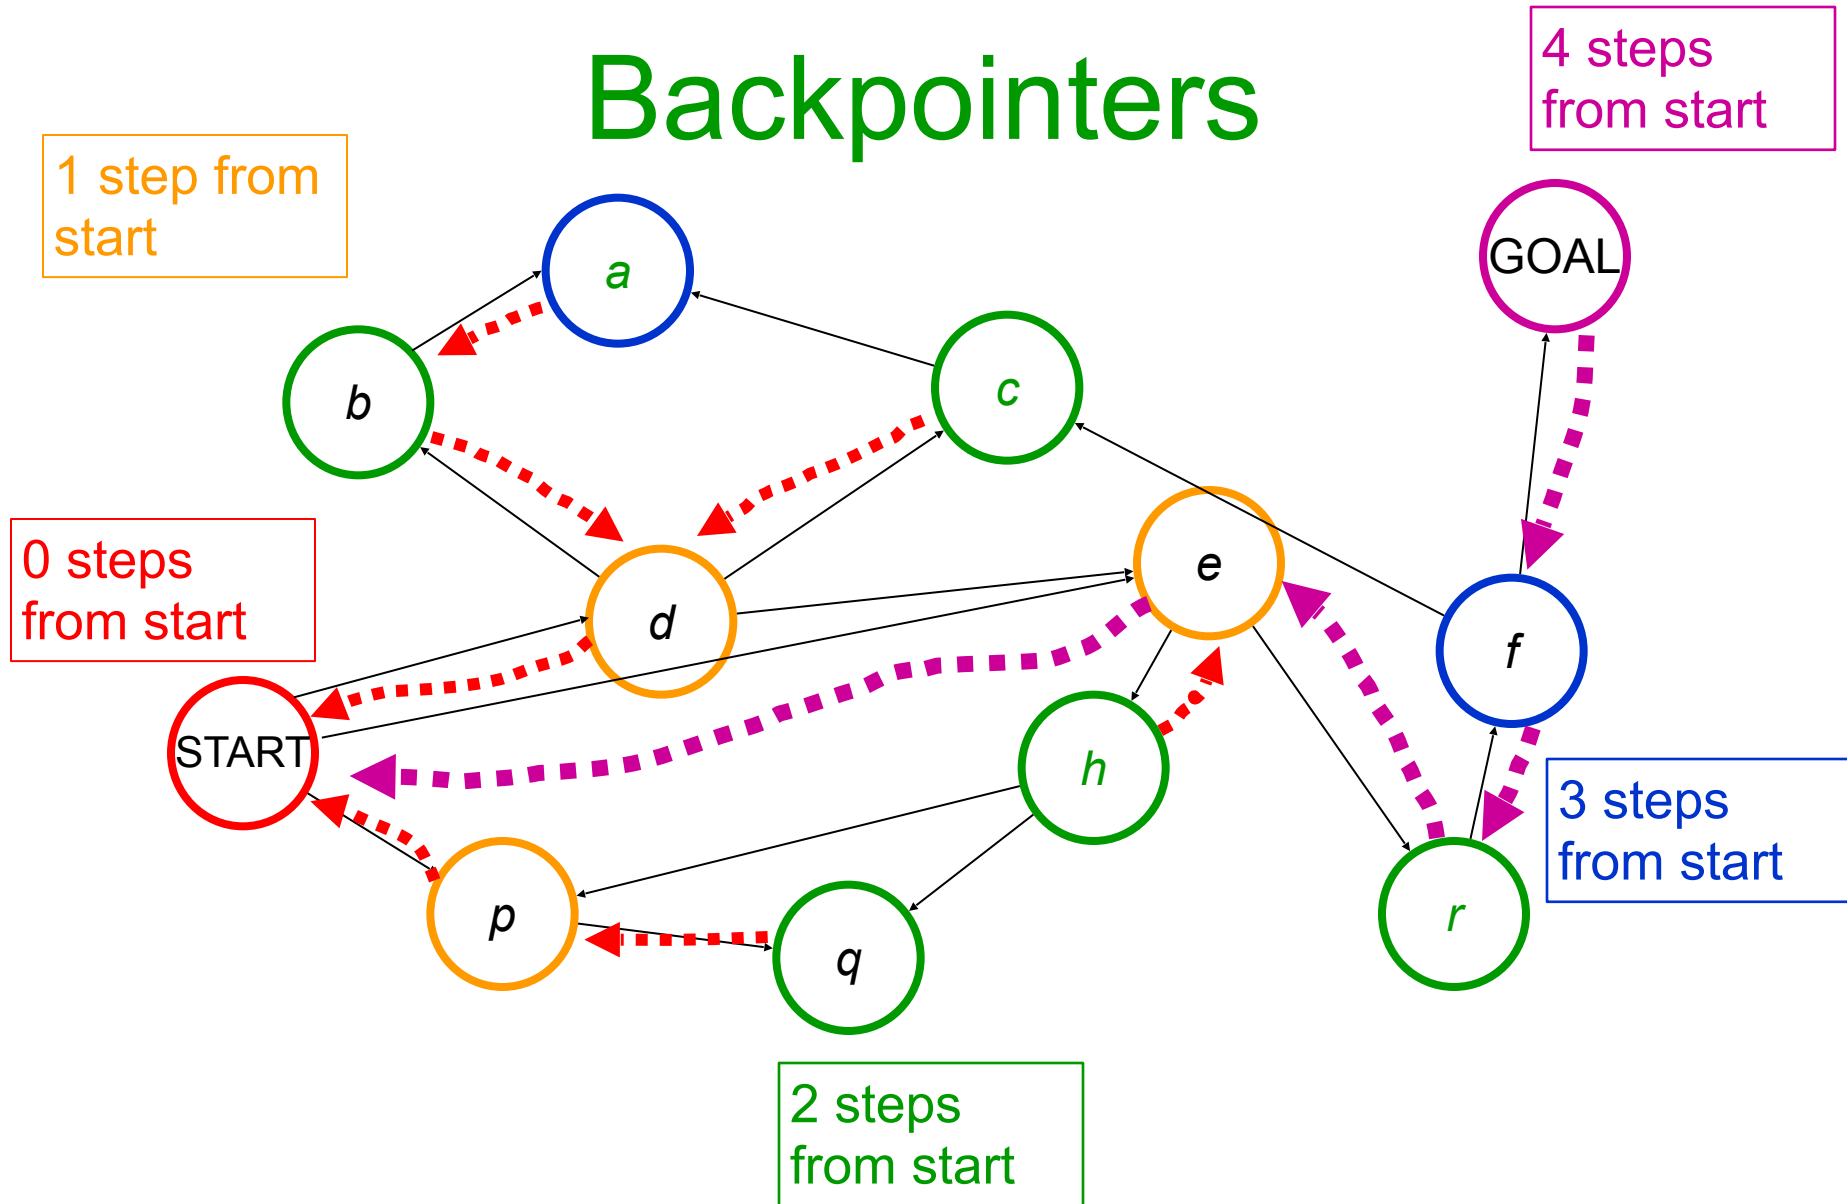
"I've got to the goal. I see I was at *f* before this. And I was at *r* before I was at *f*. And I was...

.... so solution path is $S \rightarrow e \rightarrow r \rightarrow f \rightarrow G$ "

Backpointers



Backpointers



Starting Breadth First Search

For any state s that we've labeled, we'll remember:

- $previous(s)$ as the previous state on a shortest path from START state to s .


On the k th iteration of the algorithm we'll begin with V_k defined as the set of those states for which the shortest path from the start costs exactly k steps

Then, during that iteration, we'll compute V_{k+1} , defined as the set of those states for which the shortest path from the start costs exactly $k+1$ steps


We begin with $k = 0$, $V_0 = \{\text{START}\}$ and we'll define, $previous(\text{START}) = \text{NULL}$

Then we'll add in things one step from the START into V_1 . And we'll keep going.


BFS




BFS




BFS



BFS



BFS



Breadth First Search

$V_0 := S$ (the set of start states)

$previous(START) := NIL$

$k := 0$

while (no goal state is in V_k and V_k is not empty) **do**

$V_{k+1} :=$ empty set

 For each state s in V_k

 For each state s' in **succs**(s)

 If s' has not already been labeled

 Set $previous(s') := s$

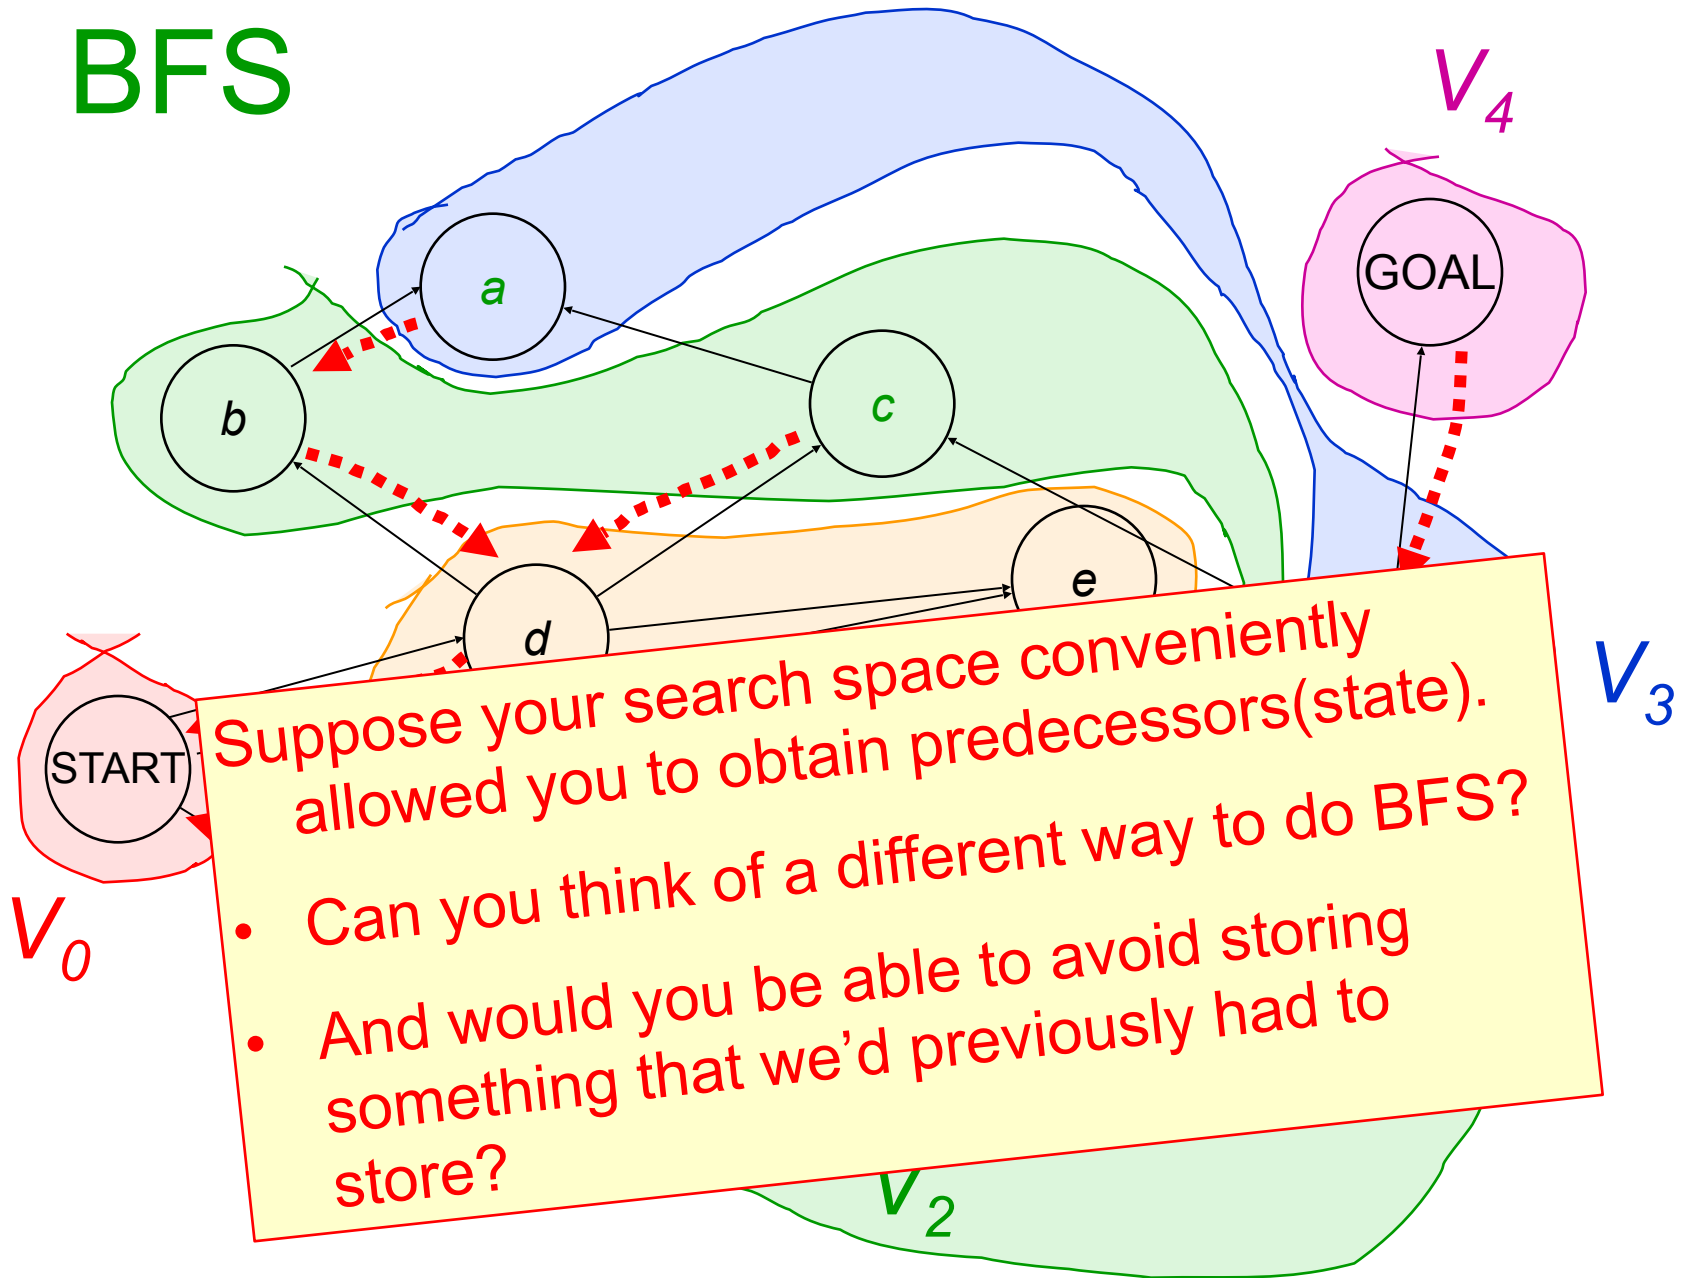
 Add s' into V_{k+1}

$k := k+1$


If V_k is empty signal FAILURE

Else build the solution path thus: Let S_i be the i th state in the shortest path. Define $S_k = GOAL$, and for all $i \leq k$, define $S_{i-1} = previous(S_i)$.

BFS



Another way: Work back



Label all states that can reach G in 1 step but can't reach it in less than 1 step.

Label all states that can reach G in 2 steps but can't reach it in less than 2 steps.


Etc. ... until start is reached.

“number of steps to goal” labels determine the shortest path. Don't need extra bookkeeping info.

Breadth First Details

- It is fine for there to be more than one goal state.
- It is fine for there to be more than one start state.
- This algorithm works forwards from the start. Any algorithm which works forwards from the start is said to be *forward chaining*.
- You can also work backwards from the goal. This algorithm is very similar to Dijkstra's algorithm.
- Any algorithm which works backwards from the goal is said to be *backward chaining*.
- Backward versus forward. Which is better?

Costs on transitions



Notice that BFS finds the shortest path in terms of number of transitions. It does not find the least-cost path.

We will quickly review an algorithm which does find the least-cost path. On the k th iteration, for any state S , write $g(s)$ as the least-cost path to S in k or fewer steps.

Least Cost Breadth First

V_k = the set of states which can be reached in exactly k steps, and for which the least-cost k -step path is less cost than any path of length less than k . In other words, V_k = the set of states whose values changed on the previous iteration.

$V_0 := S$ (the set of start states)

$previous(START) := NIL$

$g(START) = 0$

$k := 0$

while (V_k is not empty) **do**

$V_{k+1} :=$ empty set

 For each state s in V_k

 For each state s' in **succs**(s)

 If s' has not already been labeled

 OR if $g(s) + Cost(s, s') < g(s')$

 Set $previous(s') := s$

 Set $g(s') := g(s) + Cost(s, s')$

 Add s' into V_{k+1}

$k := k+1$

If GOAL not labeled, exit signaling FAILURE

Else build the solution path thus: Let S_k be the k th state in the shortest path. Define $S_k = GOAL$, and forall $i \leq k$, define $S_{i-1} = previous(S_i)$.

Uniform-Cost Search

- A conceptually simple BFS approach when there are costs on transitions
- It uses priority queues



Priority Queue Refresher

A priority queue is a data structure in which you can insert and retrieve (thing, value) pairs with the following operations:

Init-PriQueue(PQ)	initializes the PQ to be empty.
Insert-PriQueue(PQ, thing, value)	inserts <i>(thing, value)</i> into the queue.
Pop-least(PQ)	returns the <i>(thing, value)</i> pair with the lowest value, and removes it from the queue.



Priority Queue Refresher

A priority queue is a data structure in which you can insert and retrieve *(thing, value)* pairs with the following operations:

For more details, see Knuth or Sedgwick or basically any book with the word “algorithms” prominently appearing in the title.

Init-PriQueue(PQ)	initializes the PQ to be empty.
Insert-PriQueue(PQ, thing, value)	inserts <i>(thing, value)</i> into the queue.
Pop-least(PQ)	returns the <i>(thing, value)</i> pair with the lowest value, and removes it from the queue.

Priority Queues can be implemented in such a way that the cost of the insert and pop operations are

Very cheap (though not absolutely, incredibly cheap!)

$O(\log(\text{number of things in priority queue}))$


Uniform-Cost Search

- A conceptually simple BFS approach when there are costs on transitions
- It uses a priority queue

PQ = Set of states that have been expanded or are awaiting expansion

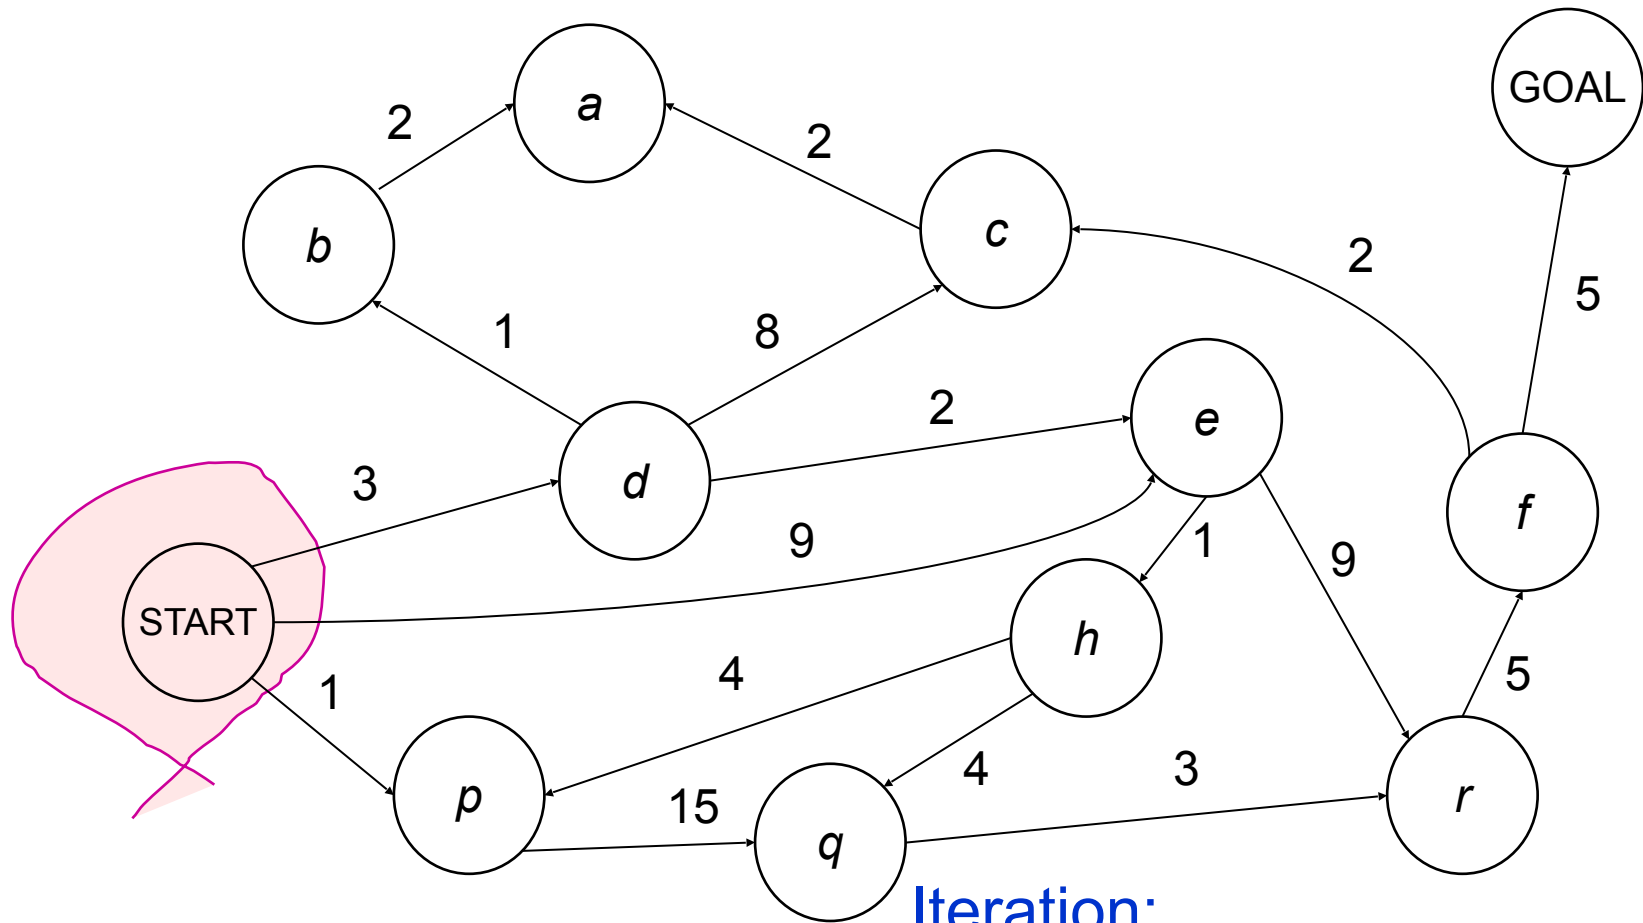
Priority of state $s = g(s)$ = cost of getting to s using path implied by backpointers.

Starting UCS



$$PQ = \{ (S, 0) \}$$

UCS Iterations

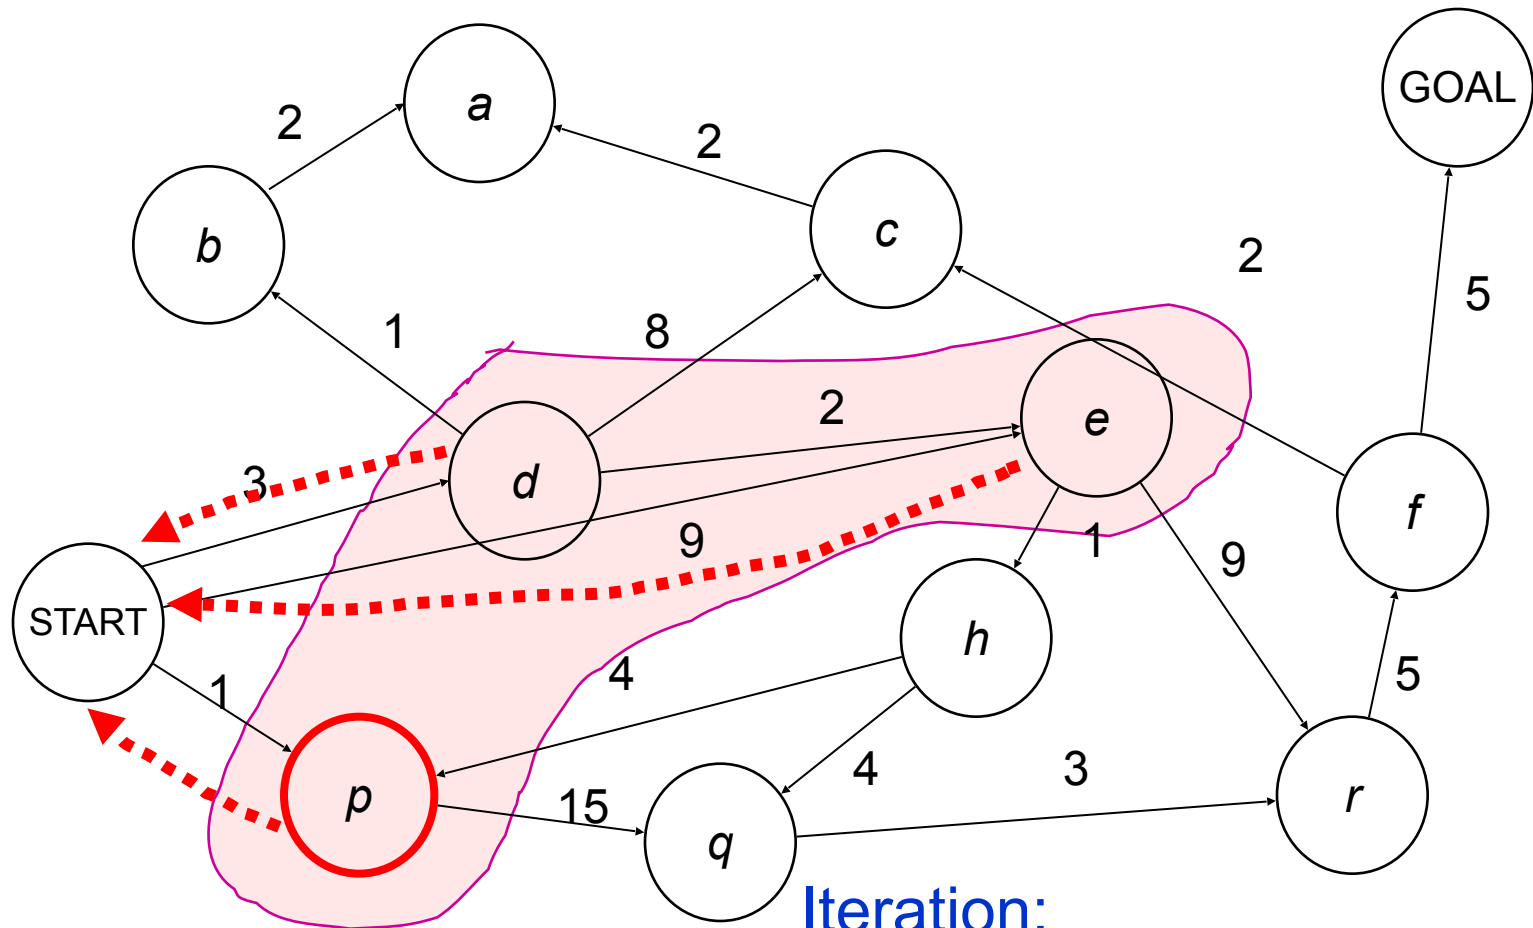


Iteration:

1. Pop least-cost state from PQ
2. Add successors

$PQ = \{ (S, 0) \}$

UCS Iterations




Iteration:

1. Pop least-cost state from PQ
2. Add successors

$PQ = \{ (p, 1), (d, 3), (e, 9) \}$


UCS Iterations



$PQ = \{ (d, 3) , (e, 9) , (q, 16) \}$


- Iteration:
1. Pop least-cost state from PQ
 2. Add successors

UCS Iterations



$PQ = \{ (b, 4) , (e, 5) , (c, 11) , (q, 16) \}$


UCS Iterations



$PQ = \{ (b, 4) , (e, 5) \}$

1. Extract least-cost state from PQ
2. Add successors

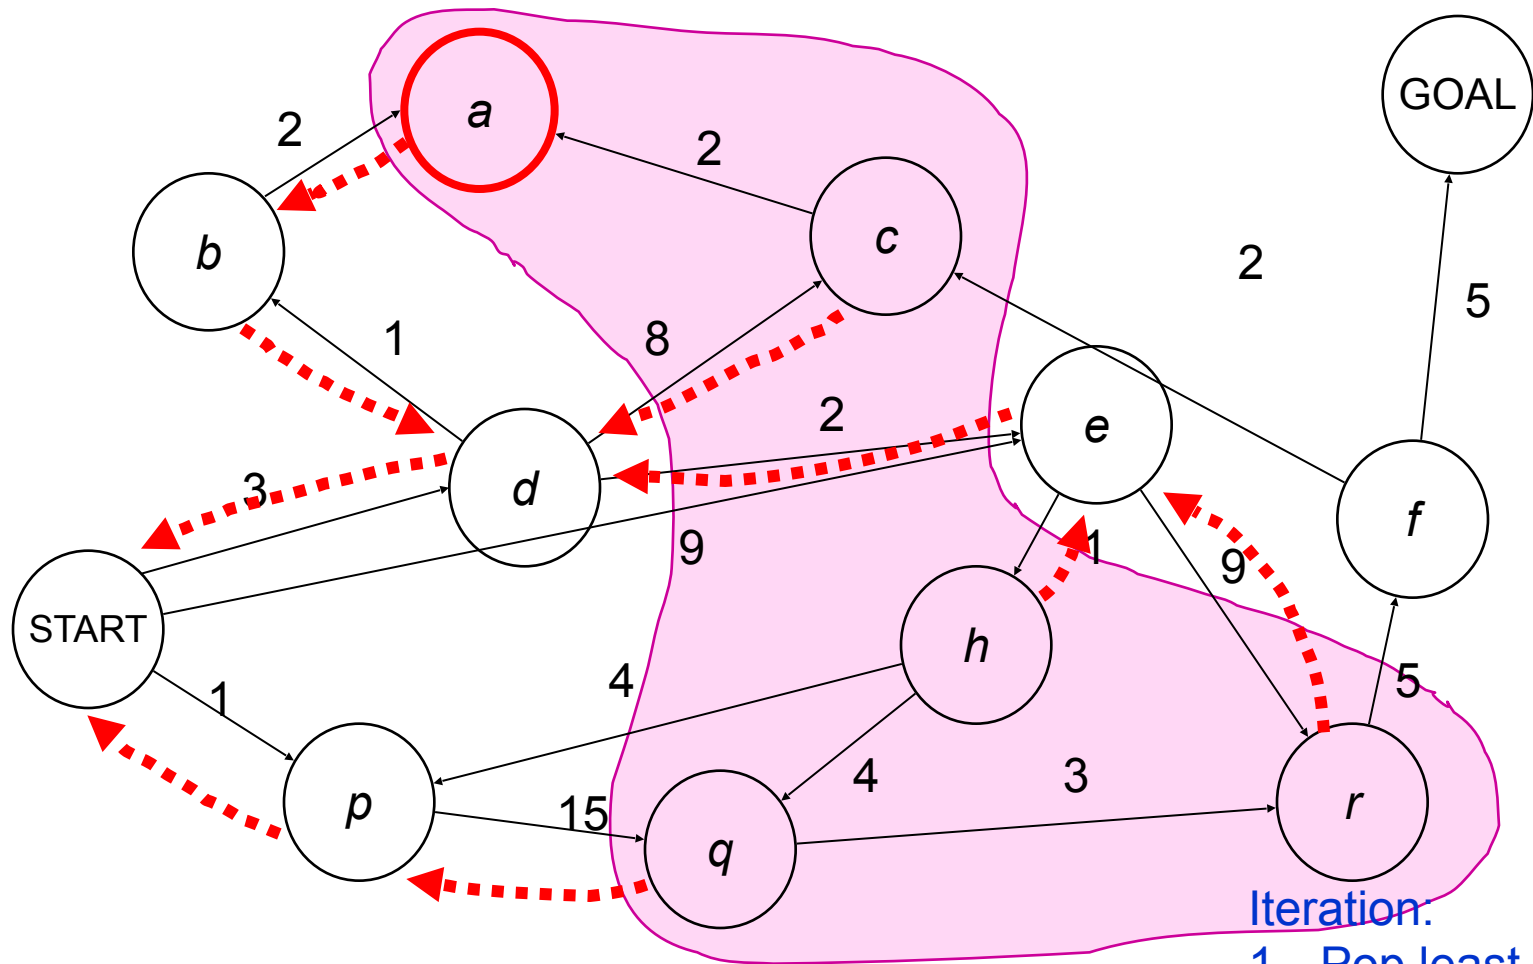
UCS Iterations



$PQ = \{ (e, 5) , (a, 6) , (c, 11) , (q, 16) \}$


- Iteration:
1. Pop least-cost state from PQ
 2. Add successors

UCS Iterations




$PQ = \{ (a, 6), (h, 6), (c, 11), (r, 14), (q, 16) \}$

UCS Iterations



$PQ = \{ (h, 6), (c, 11), (r, 14), (q, 16) \}$

UCS Iterations




$PQ = \{ (q, 10), (c, 11), (r, 14) \}$

- Iteration:
1. Pop least-cost state from PQ
 2. Add successors

Note what happened here:

- ***h*** found a new way to get to ***p***
- but it was more costly than the best known way
- and so ***p***'s priority was unchanged




Iteration:

1. Pop least-cost state from PQ
2. Add successors

$$PQ = \{ (q, 10), (c, 11), (r, 14) \}$$


UCS Iterations



$PQ = \{ (c, 11), (r, 13) \}$

- Iteration:
1. Pop least-cost state from PQ
 2. Add successors


UCS Iterations



$PQ = \{ (r, 13) \}$

- Iteration:
1. Pop least-cost state from PQ
 2. Add successors


UCS Iterations



$PQ = \{ (f, 18) \}$

- Iteration:
1. Pop least-cost state from PQ
 2. Add successors

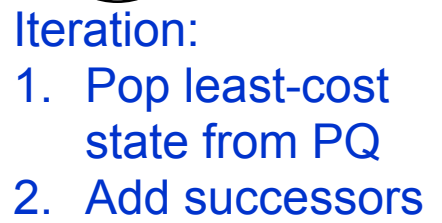
UCS Iterations




$PQ = \{ (G, 23) \}$

- Iteration:
1. Pop least-cost state from PQ
 2. Add successors

How to


$$PQ = \{ (G, 23) \}$$

UCS terminates



Judging a search algorithm

- **Completeness**: is the algorithm guaranteed to find a solution if a solution exists?
- Guaranteed to find **optimal?** (will it find the least cost path?)
- Algorithmic **time complexity**
- **Space complexity** (memory use)

Variables:

N	number of states in the problem
B	the average branching factor (the average number of successors) ($B > 1$)
L	the length of the path from start to goal with the shortest number of steps

How would we judge our algorithms?

Judging a search algorithm

N	number of states in the problem
B	the average branching factor (the average number of successors) ($B > 1$)
L	the length of the path from start to goal with the shortest number of steps
Q	the average size of the priority queue



Algorithm		Complete	Optimal	Time	Space
BFS	Breadth First Search				
LCBFS	Least Cost BFS				
UCS	Uniform Cost Search				

Judging a search algorithm

N	number of states in the problem
B	the average branching factor (the average number of successors) ($B > 1$)
L	the length of the path from start to goal with the shortest number of steps
Q	the average size of the priority queue


Algorithm		Complete	Optimal	Time	Space
BFS	Breadth First Search	Y	if all transitions same cost	$O(\min(N, B^L))$	$O(\min(N, B^L))$
LCBFS	Least Cost BFS	Y	Y	$O(\min(N, B^L))$	$O(\min(N, B^L))$
UCS	Uniform Cost Search	Y	Y	$O(\log(Q) * \min(N, B^L))$	$O(\min(N, B^L))$

Search Tree Representation



What order do we go through the search tree with BFS?


Depth First Search



An alternative to BFS. Always expand from the most-recently-expanded node, if it has any untried successors. Else backup to the previous node on the current path.



DFS in action

START
START *d*
START *d b*
START *d b a*
START *d c*
START *d c a*
START *d e*
START *d e r*
START *d e r f*
START *d e r f c*
START *d e r f c a*
START *d e r f* GOAL



DFS Search tree traversal

Can you draw in the order in which the search-tree nodes are visited?




DFS Algorithm

We use a data structure we'll call a Path to represent the , er, path from the START to the current state.

E.G. Path $P = \langle \text{START}, d, e, r \rangle$

Along with each node on the path, we must remember which successors we still have available to expand. E.G. at the following point, we'll have

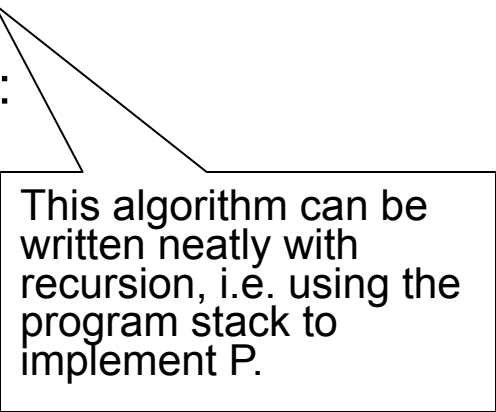


```
P = <START (expand=e , p) ,  
    d (expand = NULL) ,  
    e (expand = h) ,  
    r (expand = f) >
```

DFS Algorithm

```
Let P = <START (expand = succs(START))>
While (P not empty and top(P) not a goal)
    if expand of top(P) is empty
    then
        remove top(P) ("pop the stack")
    else
        let s be a member of expand of top(P)
        remove s from expand of top(P)
        make a new item on the top of path P:
            s (expand = succs(s))

If P is empty
    return FAILURE
Else
    return the path consisting of states in P
```



This algorithm can be written neatly with recursion, i.e. using the program stack to implement P.

Judging a search algorithm

N	number of states in the problem
B	the average branching factor (the average number of successors) ($B > 1$)
L	the length of the path from start to goal with the shortest number of steps
Q	the average size of the priority queue

Algorithm		Complete	Optimal	Time	Space
BFS	Breadth First Search	Y	if all transitions same cost	$O(\min(N, B^L))$	$O(\min(N, B^L))$
LCBFS	Least Cost BFS	Y	Y	$O(\min(N, B^L))$	$O(\min(N, B^L))$
UCS	Uniform Cost Search	Y	Y	$O(\log(Q) * \min(N, B^L))$	$O(\min(N, B^L))$
DFS	Depth First Search				

Judging a search algorithm

N	number of states in the problem
B	the average branching factor (the average number of successors) ($B > 1$)
L	the length of the path from start to goal with the shortest number of steps
Q	the average size of the priority queue

Algorithm		Complete	Optimal	Time	Space
BFS	Breadth First Search	Y	if all transitions same cost	$O(\min(N, B^L))$	$O(\min(N, B^L))$
LCBFS	Least Cost BFS	Y	Y	$O(\min(N, B^L))$	$O(\min(N, B^L))$
UCS	Uniform Cost Search	Y	Y	$O(\log(Q) * \min(N, B^L))$	$O(\min(N, B^L))$
DFS	Depth First Search	N	N	N/A	N/A

Judging a search algorithm

N	number of states in the problem
B	the average branching factor (the average number of successors) ($B > 1$)
L	the length of the path from start to goal with the shortest number of steps
Q	the average size of the priority queue

Algorithm		Complete	Optimal	Time	Space
BFS	Breadth First Search	Y	if all transitions same cost	$O(\min(N, B^L))$	$O(\min(N, B^L))$
LCBFS	Least Cost BFS	Y	Y	$O(\min(N, B^L))$	$O(\min(N, B^L))$
UCS	Uniform Cost Search	Y	Y	$O(\log(Q) * \min(N, B^L))$	$O(\min(N, B^L))$
DFS**	Depth First Search				

Assuming Acyclic Search Space

Judging a search algorithm

N	number of states in the problem
B	the average branching factor (the average number of successors) ($B > 1$)
L	the length of the path from start to goal with the shortest number of steps
LMAX	Length of longest path from start to anywhere
Q	the average size of the priority queue

Algorithm		Complete	Optimal	Time	Space
BFS	Breadth First Search	Y	if all transitions same cost	$O(\min(N, B^L))$	$O(\min(N, B^L))$
LCBFS	Least Cost BFS	Y	Y	$O(\min(N, B^L))$	$O(\min(N, B^L))$
UCS	Uniform Cost Search	Y	Y	$O(\log(Q) * \min(N, B^L))$	$O(\min(N, B^L))$
DFS**	Depth First Search	Y	N	$O(B^{LMAX})$	$O(LMAX)$

Assuming Acyclic Search Space

Questions to ponder

- How would you prevent DFS from looping?
- How could you force it to give an optimal solution?

Questions to ponder

- How would you prevent DFS from looping?
- How could you force it to give an optimal solution?

Answer 1:

PC-DFS (Path Checking DFS):

Answer 2:

MEMDFS (Memoizing DFS):

Questions to ponder

- How would you prevent DFS from looping?
- How could you force it to give an optimal solution?

Answer 1:

PC-DFS (Path Checking DFS):

Don't recurse on a state if that state is already in the current path

Answer 2:

MEMDFS (Memoizing DFS):

Remember all states expanded so far. Never expand anything twice.

Questions to ponder

- How would you prevent DFS from looping?

Answer 1:

PC-DFS (Path Checking DFS):

Don't recurse on a state if that state is already in the current path

Answer 2:

MEMDFS (Memoizing DFS):

Remember all states expanded so far. Never expand anything twice.

Are there occasions when PCDFS is better than MEMDFS?

Are there occasions when MEMDFS is better than PCDFS?

How can we give an

optimal solution?

Judging a search algorithm

N	number of states in the problem
B	the average branching factor (the average number of successors) ($B > 1$)
L	the length of the path from start to goal with the shortest number of steps
LMAX	Length of longest cycle-free path from start to anywhere
Q	the average size of the priority queue

Algorithm		Complete	Optimal	Time	Space
BFS	Breadth First Search	Y	if all transitions same cost	$O(\min(N, B^L))$	$O(\min(N, B^L))$
LCBFS	Least Cost BFS	Y	Y	$O(\min(N, B^L))$	$O(\min(N, B^L))$
UCS	Uniform Cost Search	Y	Y	$O(\log(Q) * \min(N, B^L))$	$O(\min(N, B^L))$
PCDFS	Path Check DFS				
MEMDFS	Memoizing DFS				

Judging a search algorithm

N	number of states in the problem
B	the average branching factor (the average number of successors) ($B > 1$)
L	the length of the path from start to goal with the shortest number of steps
LMAX	Length of longest cycle-free path from start to anywhere
Q	the average size of the priority queue

Algorithm		Complete	Optimal	Time	Space
BFS	Breadth First Search	Y	if all transitions same cost	$O(\min(N, B^L))$	$O(\min(N, B^L))$
LCBFS	Least Cost BFS	Y	Y	$O(\min(N, B^L))$	$O(\min(N, B^L))$
UCS	Uniform Cost Search	Y	Y	$O(\log(Q) * \min(N, B^L))$	$O(\min(N, B^L))$
PCDFS	Path Check DFS	Y	N	$O(B^{LMAX})$	$O(LMAX)$
MEMDFS	Memoizing DFS	Y	N	$O(\min(N, B^{LMAX}))$	$O(\min(N, B^{LMAX}))$


Judging a search algorithm

N	number of states in the problem
B	the average branching factor (the average number of successors) ($B > 1$)
L	the length of the path from start to goal with the shortest number of steps
LMAX	Length of longest cycle-free path from start to anywhere
Q	the average size of the priority queue

Algorithm		Complete	Optimal	Time	Space
BFS	Breadth First Search	Y	if all transitions same cost	$O(\min(N, B^L))$	$O(\min(N, B^L))$
LCBFS	Least Cost BFS	Y	Y	$O(\min(N, B^L))$	$O(\min(N, B^L))$
UCS	Uniform Cost Search	Y	Y	$O(\log(Q) * \min(N, B^L))$	$O(\min(N, B^L))$
PCDFS	Path Check DFS	Y	N	$O(B^{LMAX})$	$O(LMAX)$
MEMDFS	Memoizing DFS	Y	N	$O(\min(N, B^{LMAX}))$	$O(\min(N, B^{LMAX}))$

Maze example

Imagine states are cells in a maze, you can move N, E, S, W. What would **plain DFS** do, assuming it always expanded the E successor first, then N, then W, then S?




Other questions:

What would BFS do?


What would PCDFS do?

What would MEMDFS do?

Two other DFS examples



Order: N, E, S, W?



Order: N, E, S, W
with loops prevented

Forward DFSearch or Backward DFSearch

If you have a predecessors() function as well as a successors() function you can begin at the goal and depth-first-search backwards until you hit a start.

Why/When might this be a good idea?

Invent An Algorithm Time!

Here's a way to dramatically decrease costs sometimes. Bidirectional Search. Can you guess what this algorithm is, and why it can be a huge cost-saver?

N	number of states in the problem
B	the average branching factor (the average number of successors) ($B > 1$)
L	the length of the path from start to goal with the shortest number of steps
LMAX	Length of longest cycle-free path from start to anywhere
Q	the average size of the priority queue

Algorithm		Complete	Optimal	Time	Space
BFS	Breadth First Search	Y	if all transitions same cost	$O(\min(N, B^L))$	$O(\min(N, B^L))$
LCBFS	Least Cost BFS	Y	Y	$O(\min(N, B^L))$	$O(\min(N, B^L))$
UCS	Uniform Cost Search	Y	Y	$O(\log(Q) * \min(N, B^L))$	$O(\min(N, B^L))$
PCDFS	Path Check DFS	Y	N	$O(B^{LMAX})$	$O(LMAX)$
MEMDFS	Memoizing DFS	Y	N	$O(\min(N, B^{LMAX}))$	$O(\min(N, B^{LMAX}))$
BIBFS	Bidirectional BF Search				

N	number of states in the problem
B	the average branching factor (the average number of successors) ($B > 1$)
L	the length of the path from start to goal with the shortest number of steps
LMAX	Length of longest cycle-free path from start to anywhere
Q	the average size of the priority queue

Algorithm		Complete	Optimal	Time	Space
BFS	Breadth First Search	Y	if all transitions same cost	$O(\min(N, B^L))$	$O(\min(N, B^L))$
LCBFS	Least Cost BFS	Y	Y	$O(\min(N, B^L))$	$O(\min(N, B^L))$
UCS	Uniform Cost Search	Y	Y	$O(\log(Q) * \min(N, B^L))$	$O(\min(N, B^L))$
PCDFS	Path Check DFS	Y	N	$O(B^{LMAX})$	$O(LMAX)$
MEMDFS	Memoizing DFS	Y	N	$O(\min(N, B^{LMAX}))$	$O(\min(N, B^{LMAX}))$
BIBFS	Bidirectional Search	Y	All trans same cost	$O(\min(N, 2B^{L/2}))$	$O(\min(N, 2B^{L/2}))$

Iterative Deepening

Iterative deepening is a simple algorithm which uses DFS as a subroutine:

1. Do a DFS which only searches for paths of length 1 or less. (DFS gives up any path of length 2)
2. If “1” failed, do a DFS which only searches paths of length 2 or less.
3. If “2” failed, do a DFS which only searches paths of length 3 or less.
....and so on until success


Cost is

$$O(b^1 + b^2 + b^3 + b^4 \dots + b^L) = O(b^L)$$

Can be much better than regular DFS. But cost can be much greater than the number of states.

Maze example

Imagine states are cells in a maze, you can move N, E, S, W. What would **Iterative Deepening** do, assuming it always expanded the E successor first, then N, then W, then S?



Expansion order E, N, W, S

N	number of states in the problem
B	the average branching factor (the average number of successors) ($B > 1$)
L	the length of the path from start to goal with the shortest number of steps
LMAX	Length of longest cycle-free path from start to anywhere
Q	the average size of the priority queue

Algorithm		Complete	Optimal	Time	Space
BFS	Breadth First Search	Y	if all transitions same cost	$O(\min(N, B^L))$	$O(\min(N, B^L))$
LCBFS	Least Cost BFS	Y	Y	$O(\min(N, B^L))$	$O(\min(N, B^L))$
UCS	Uniform Cost Search	Y	Y	$O(\log(Q) * \min(N, B^L))$	$O(\min(N, B^L))$
PCDFS	Path Check DFS	Y	N	$O(B^{LMAX})$	$O(LMAX)$
MEMDFS	Memoizing DFS	Y	N	$O(\min(N, B^{LMAX}))$	$O(\min(N, B^{LMAX}))$
BIBFS	Bidirectional Search	Y	All trans same cost	$O(\min(N, 2B^{L/2}))$	$O(\min(N, 2B^{L/2}))$
ID	Iterative Deepening				

N	number of states in the problem
B	the average branching factor (the average number of successors) ($B > 1$)
L	the length of the path from start to goal with the shortest number of steps
LMAX	Length of longest cycle-free path from start to anywhere
Q	the average size of the priority queue

Algorithm		Complete	Optimal	Time	Space
BFS	Breadth First Search	Y	if all transitions same cost	$O(\min(N, B^L))$	$O(\min(N, B^L))$
LCBFS	Least Cost BFS	Y	Y	$O(\min(N, B^L))$	$O(\min(N, B^L))$
UCS	Uniform Cost Search	Y	Y	$O(\log(Q) * \min(N, B^L))$	$O(\min(N, B^L))$
PCDFS	Path Check DFS	Y	N	$O(B^{LMAX})$	$O(LMAX)$
MEMDFS	Memoizing DFS	Y	N	$O(\min(N, B^{LMAX}))$	$O(\min(N, B^{LMAX}))$
BIBFS	Bidirectional Search	Y	All transitions same cost	$O(\min(N, 2B^{L/2}))$	$O(\min(N, 2B^{L/2}))$
ID	Iterative Deepening	Y	if all transitions same cost	$O(B^L)$	$O(L)$

Best First “Greedy” Search

Needs a *heuristic*. A heuristic function maps a state onto an estimate of the cost to the goal from that state.

Can you think of examples of heuristics?

E.G. for the 8-puzzle?

E.G. for planning a path through a maze?

Denote the heuristic by a function $h(s)$ from states to a cost value.

Heuristic Search

Suppose in addition to the standard search specification we also have a *heuristic*.

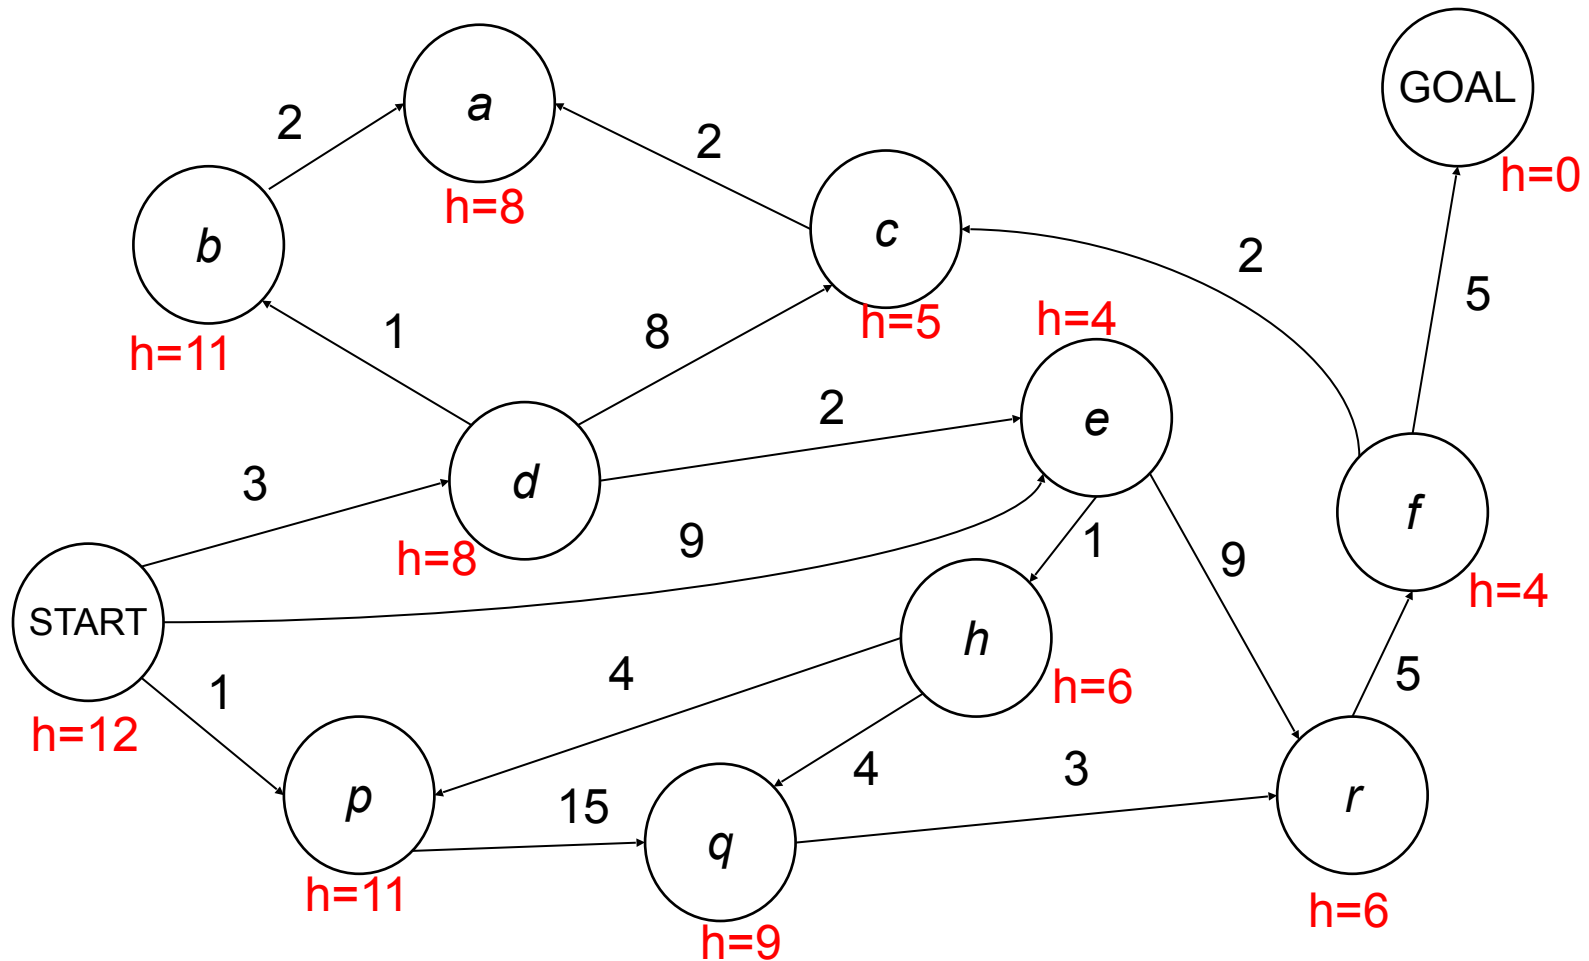
A heuristic function maps a state onto an estimate of the cost to the goal from that state.

Can you think of examples of heuristics?

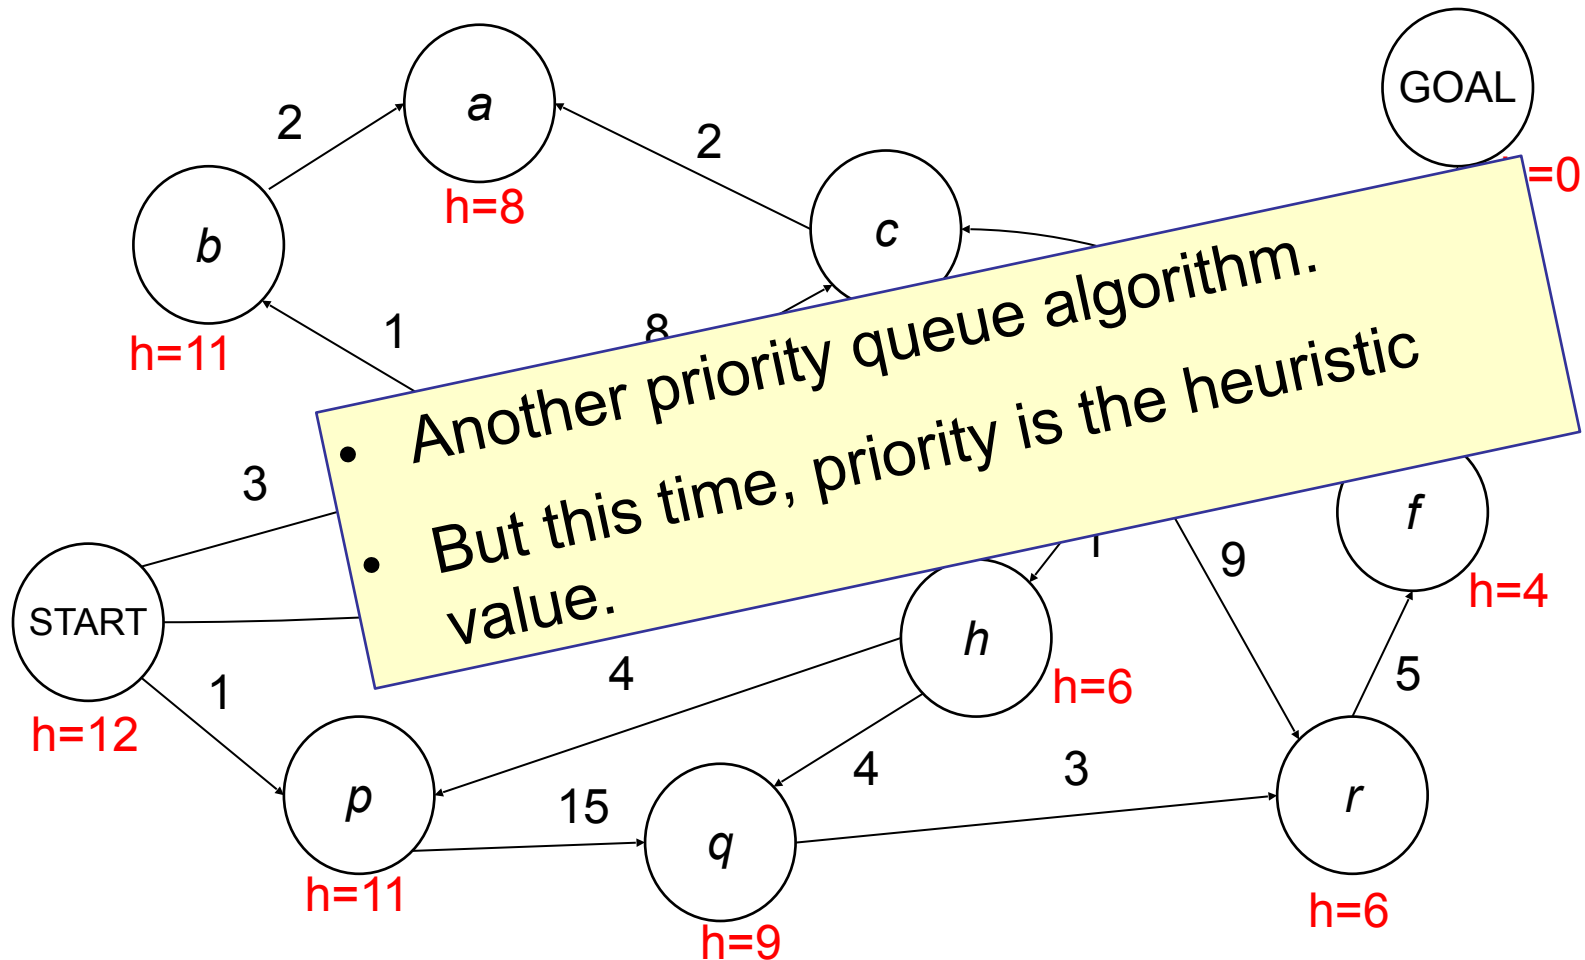
- E.G. for the 8-puzzle?
- E.G. for planning a path through a maze?

Denote the heuristic by a function $h(s)$ from states to a cost value.

Euclidian Heuristic



Euclidian Heuristic



Best First “Greedy” Search

Init-PriQueue(PQ)

Insert-PriQueue(PQ, START, h(START))

while (PQ is not empty and PQ does not contain a goal state)

 (s , h) := Pop-least(PQ)

 foreach s' in succs(s)

 if s' is not already in PQ and s' never previously been visited

 Insert-PriQueue(PQ, s', h(s'))

Algorithm		Complete	Optimal	Time	Space
BestFS	Best First Search	Y	N	$O(\min(N, B^{L_{MAX}}))$	$O(\min(N, B^{L_{MAX}}))$

better. It's a little thing we like to call: A*

...to be continued!

What you should know

- Thorough understanding of BFS, LCBFS, UCS, PCDFS, MEMDFS
- Understand the concepts of whether a search is complete, optimal, its time and space complexity
- Understand the ideas behind iterative deepening and bidirectional search
- Be able to discuss at cocktail parties the pros and cons of the above searches